

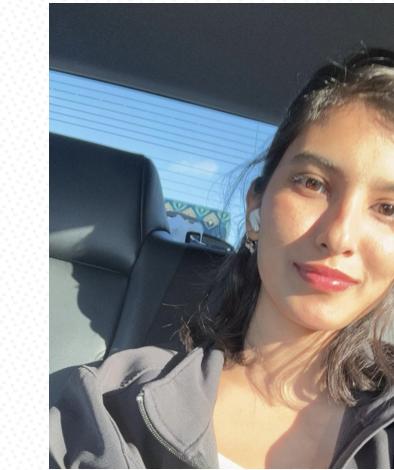
Graduates Rising in Information and Data Science

ResumeGenie: Intelligent Resume Optimization with LLMs



Aryan Vats

MS in CS



Nidhi Choudhary

MS in ADS



Eben Gunadi

MS in DS
(Healthcare)



Muqi Zhang

MS in CS



Justin Chen

BS in DS

Table of Contents



Problem Understanding

- Challenges in Traditional Resume Creation
- Need for Intelligent Automation



Solution Orchestration

- High-level Architecture
- Key Entities and Components



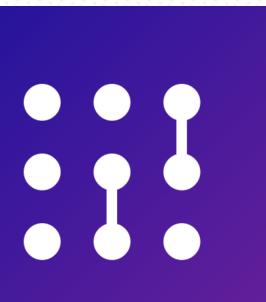
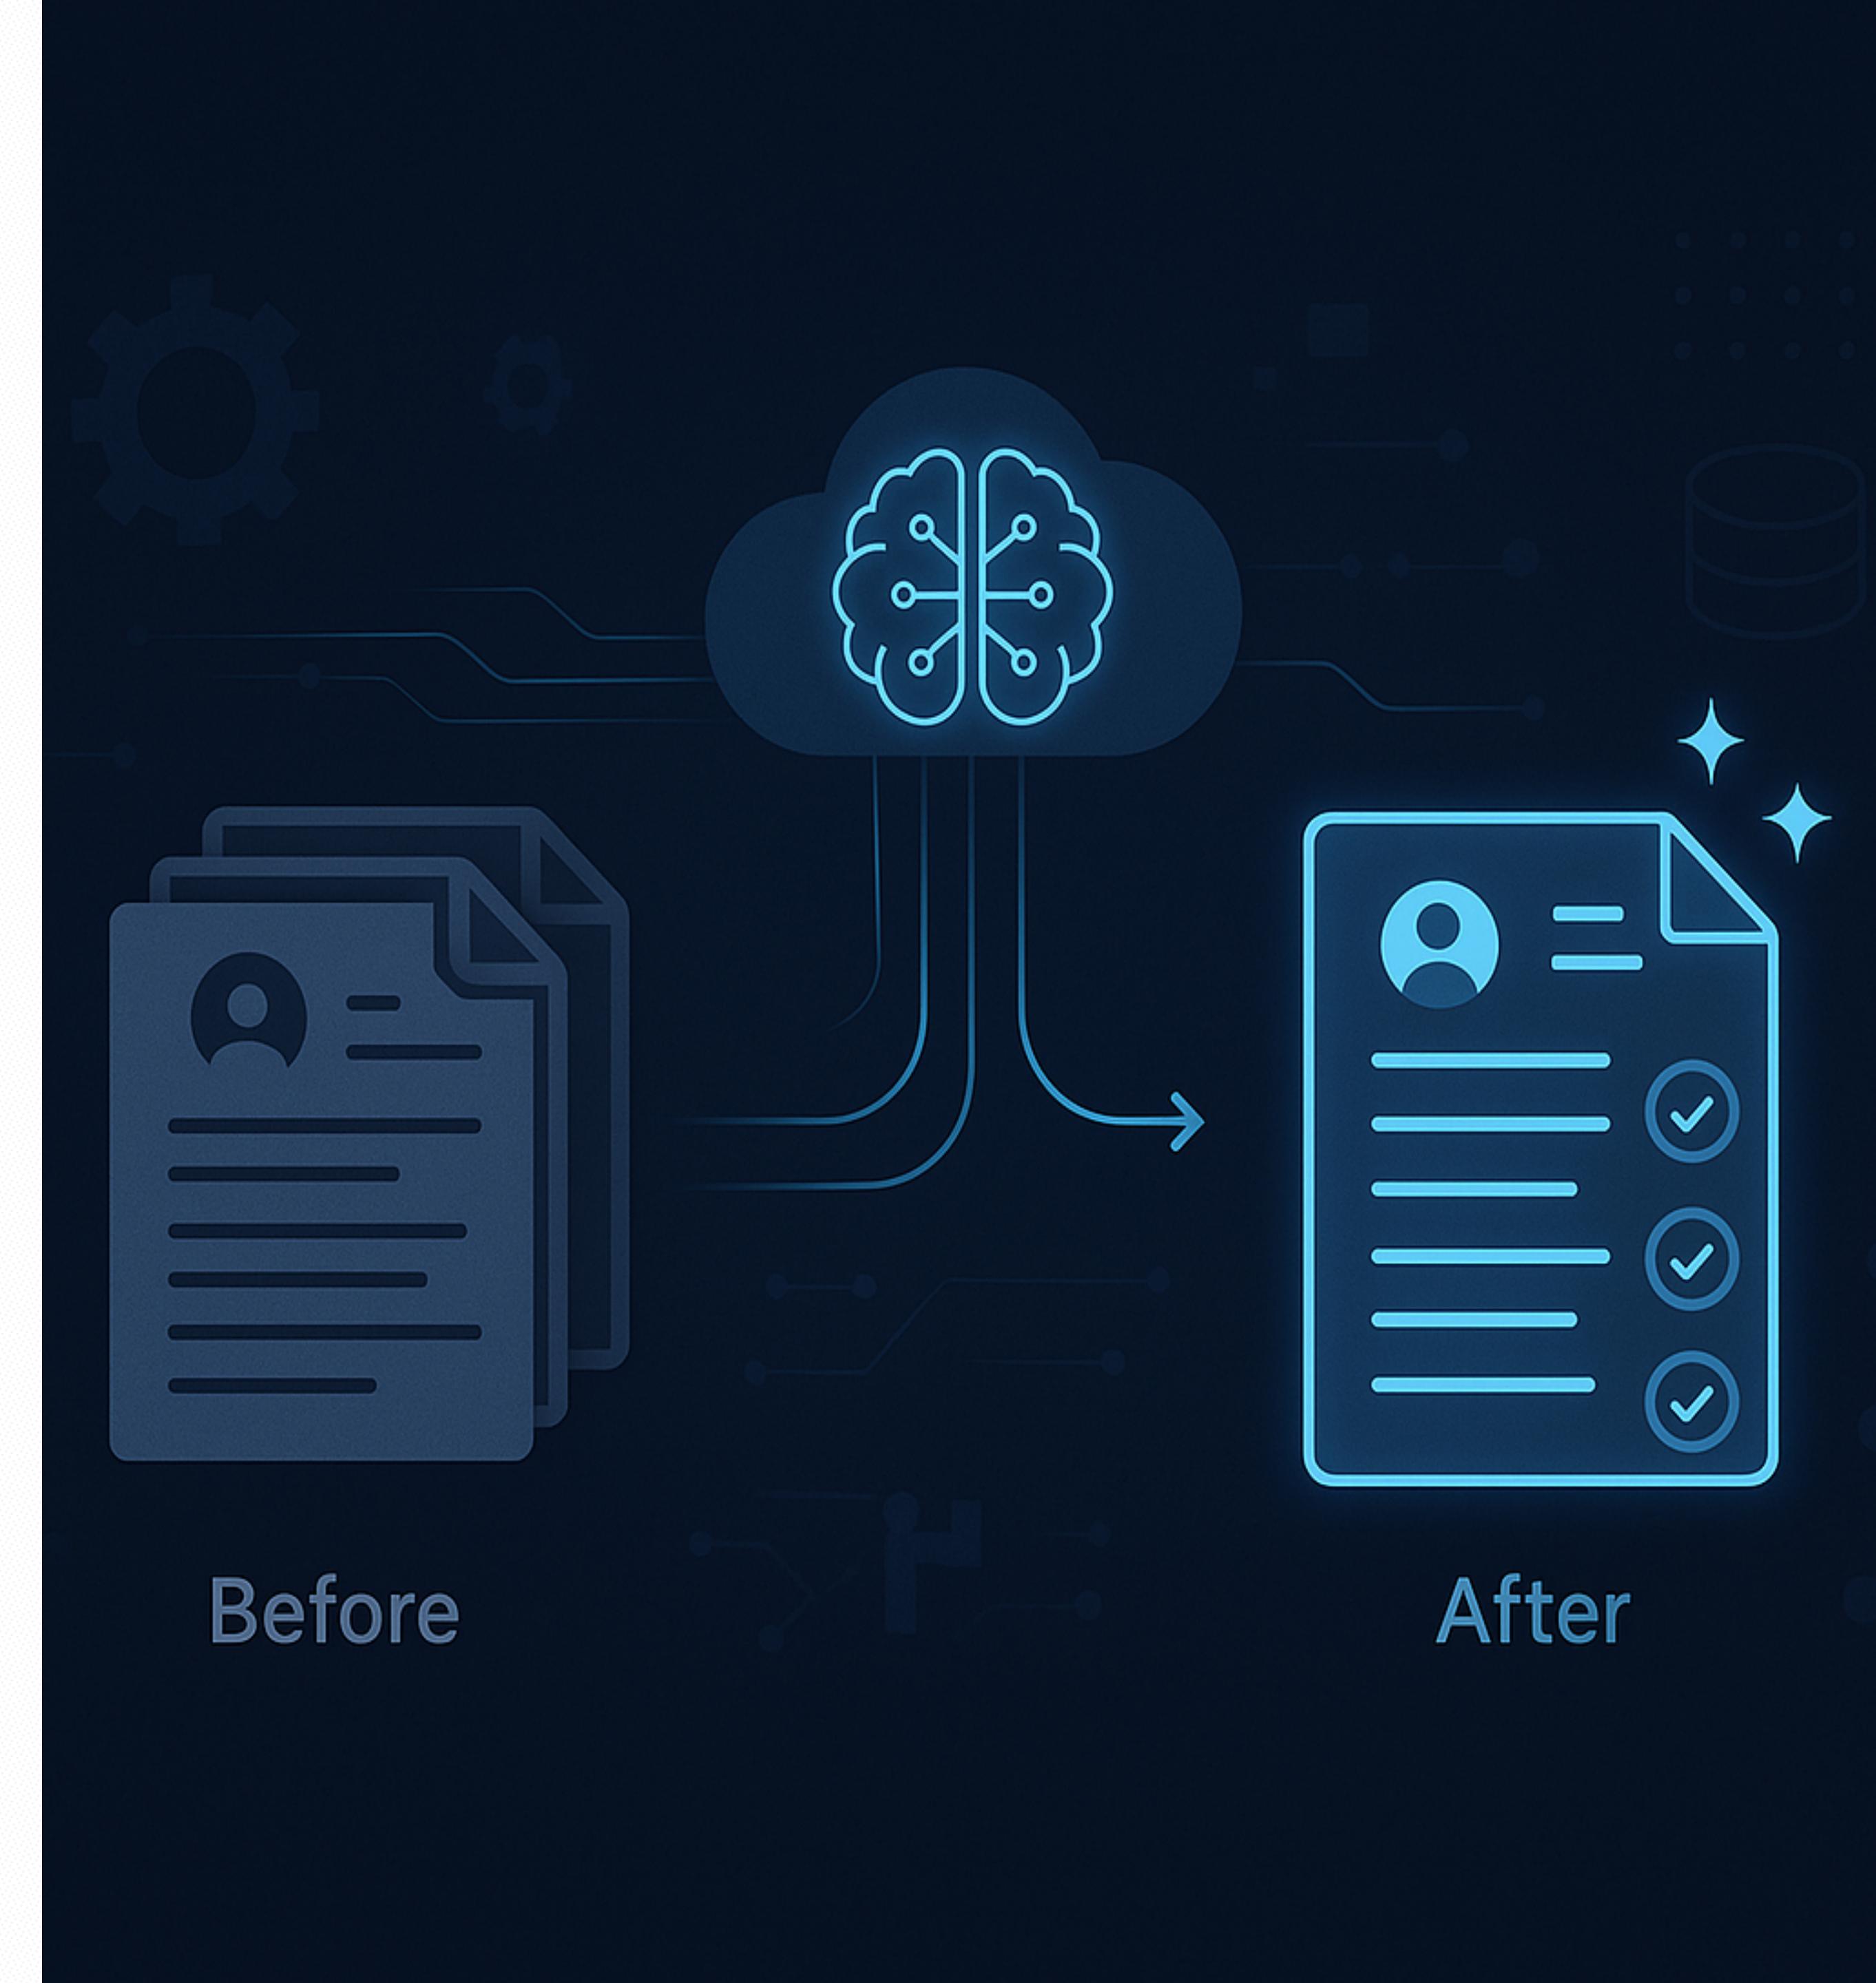
Deep Dive into the Solution Workflow

- PDF to JSON Parsing
- Prompt Engineering for Personalization
- RAG (Retrieval-Augmented Generation) for Smart Matching
- Converting JSON Back to PDF



Application Walkthrough and Result

- Previous Resume vs. New Generated Resume: Comparison and Insights



GRIDS

Why are we here?



Situation

The **job market** is highly **competitive**, and landing the right role is tougher than ever. On top of that, Applicant Tracking Systems (**ATS**) **filter out many resumes** before they even reach a recruiter.

Candidates need strong, ATS-compliant resumes to even get noticed.



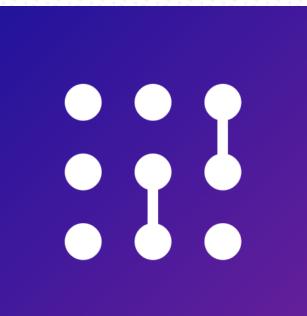
Complication

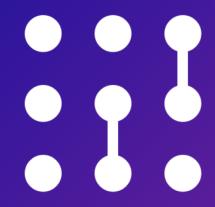
Most resumes fail to stand out due to generic content and lack of strategic positioning. Creating a **structured, impactful**, and **keyword-optimized** resume for every application is **critical** — but time-consuming. Manually tailoring resumes to each job description is unrealistic when candidates need to apply to dozens of roles quickly.



Key Question

How can we leverage technology, especially LLMs, to **automate and personalize resume generation** – making every resume tailored, ATS-friendly, and impactful – to **maximize a candidate's chances** across multiple job applications?





Graduates Rising in Information and Data Science

Solution Orchestration

Solution Orchestration

Candidate Upload: Raw input files (resume PDFs and Job Descriptions)

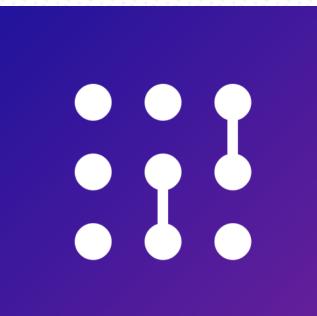
Parsing Layer: Resume parsing (PDF → JSON)

Prompt Engineering: Craft structured prompts for LLM based on parsed data

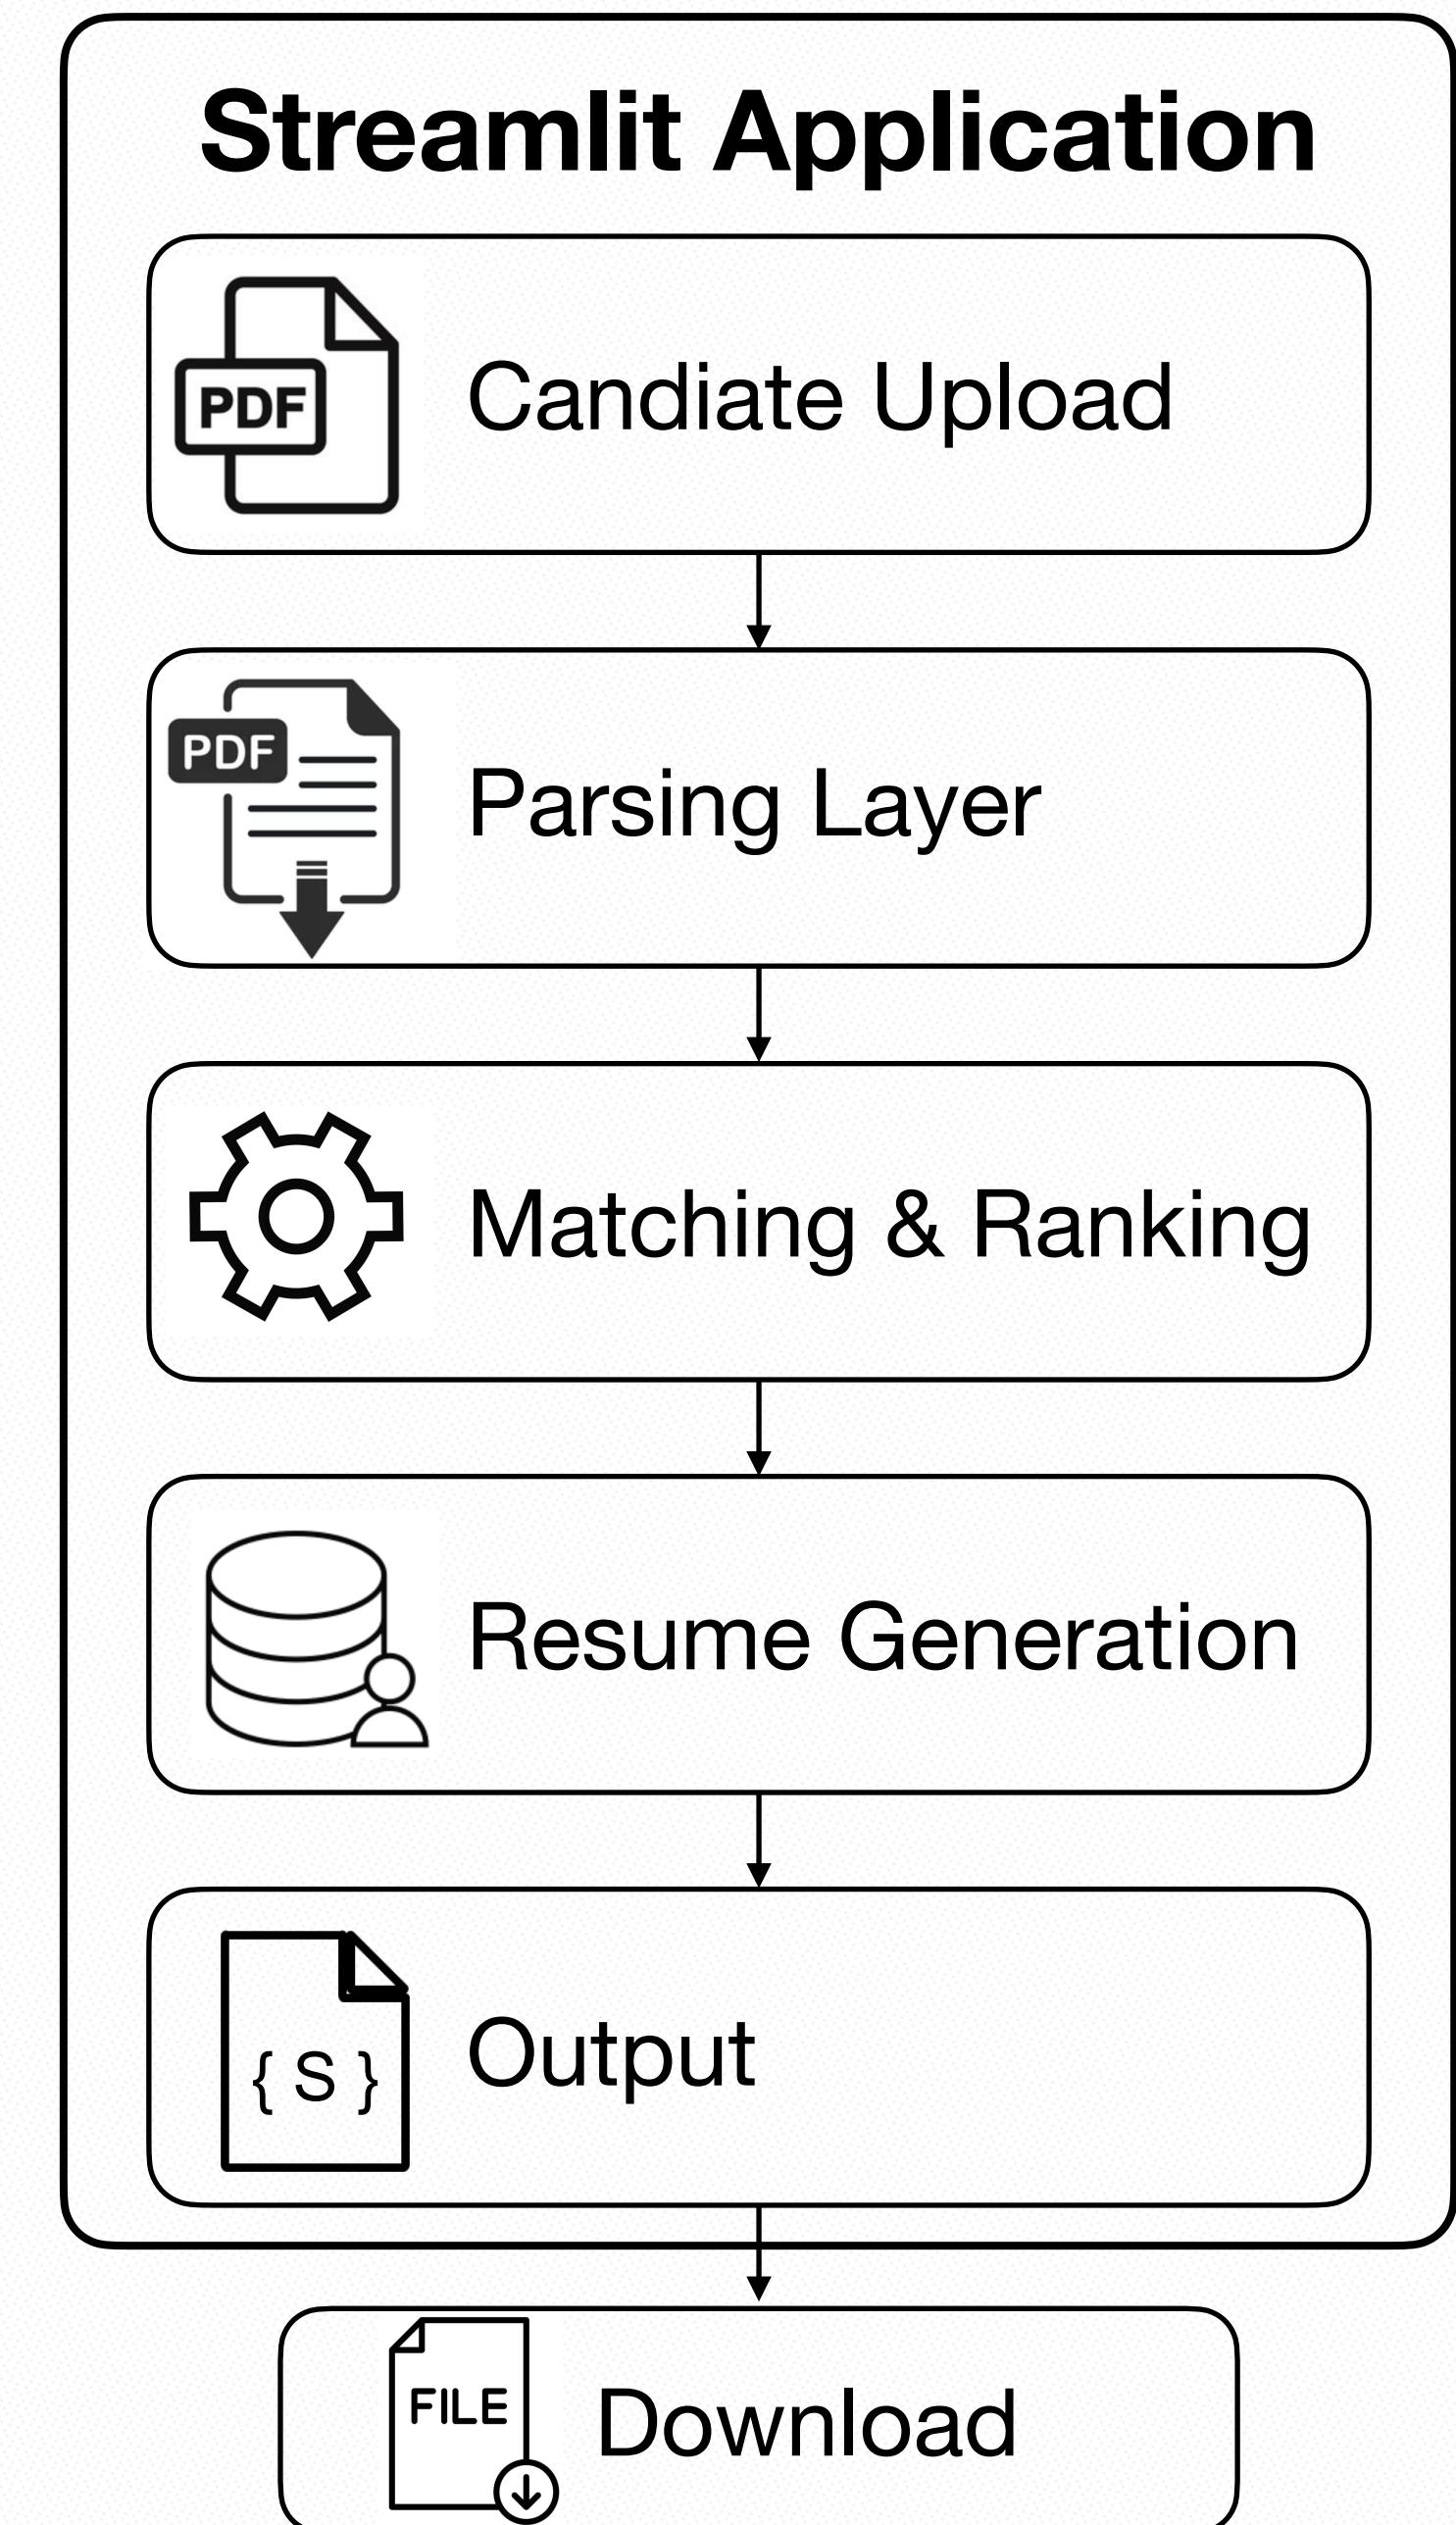
Matching & Ranking: Search and match using Vector Database (resume embeddings vs. job description embeddings)

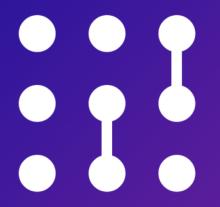
Resume Generation: Fine-tuned resume generated (using LLM with engineered prompts)

Output: Downloadable final resume (JSON → PDF)



GRIDS





Graduates Rising in Information and Data Science

Deep Dive into the Solution Workflow

Parsing with PDFPlumber, Pydantic and ChatGPT-4o

- **PDFPlumber**

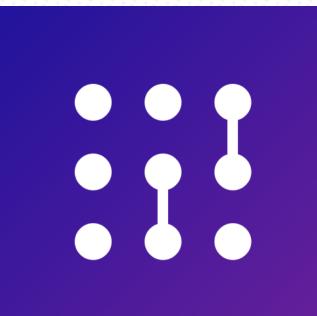
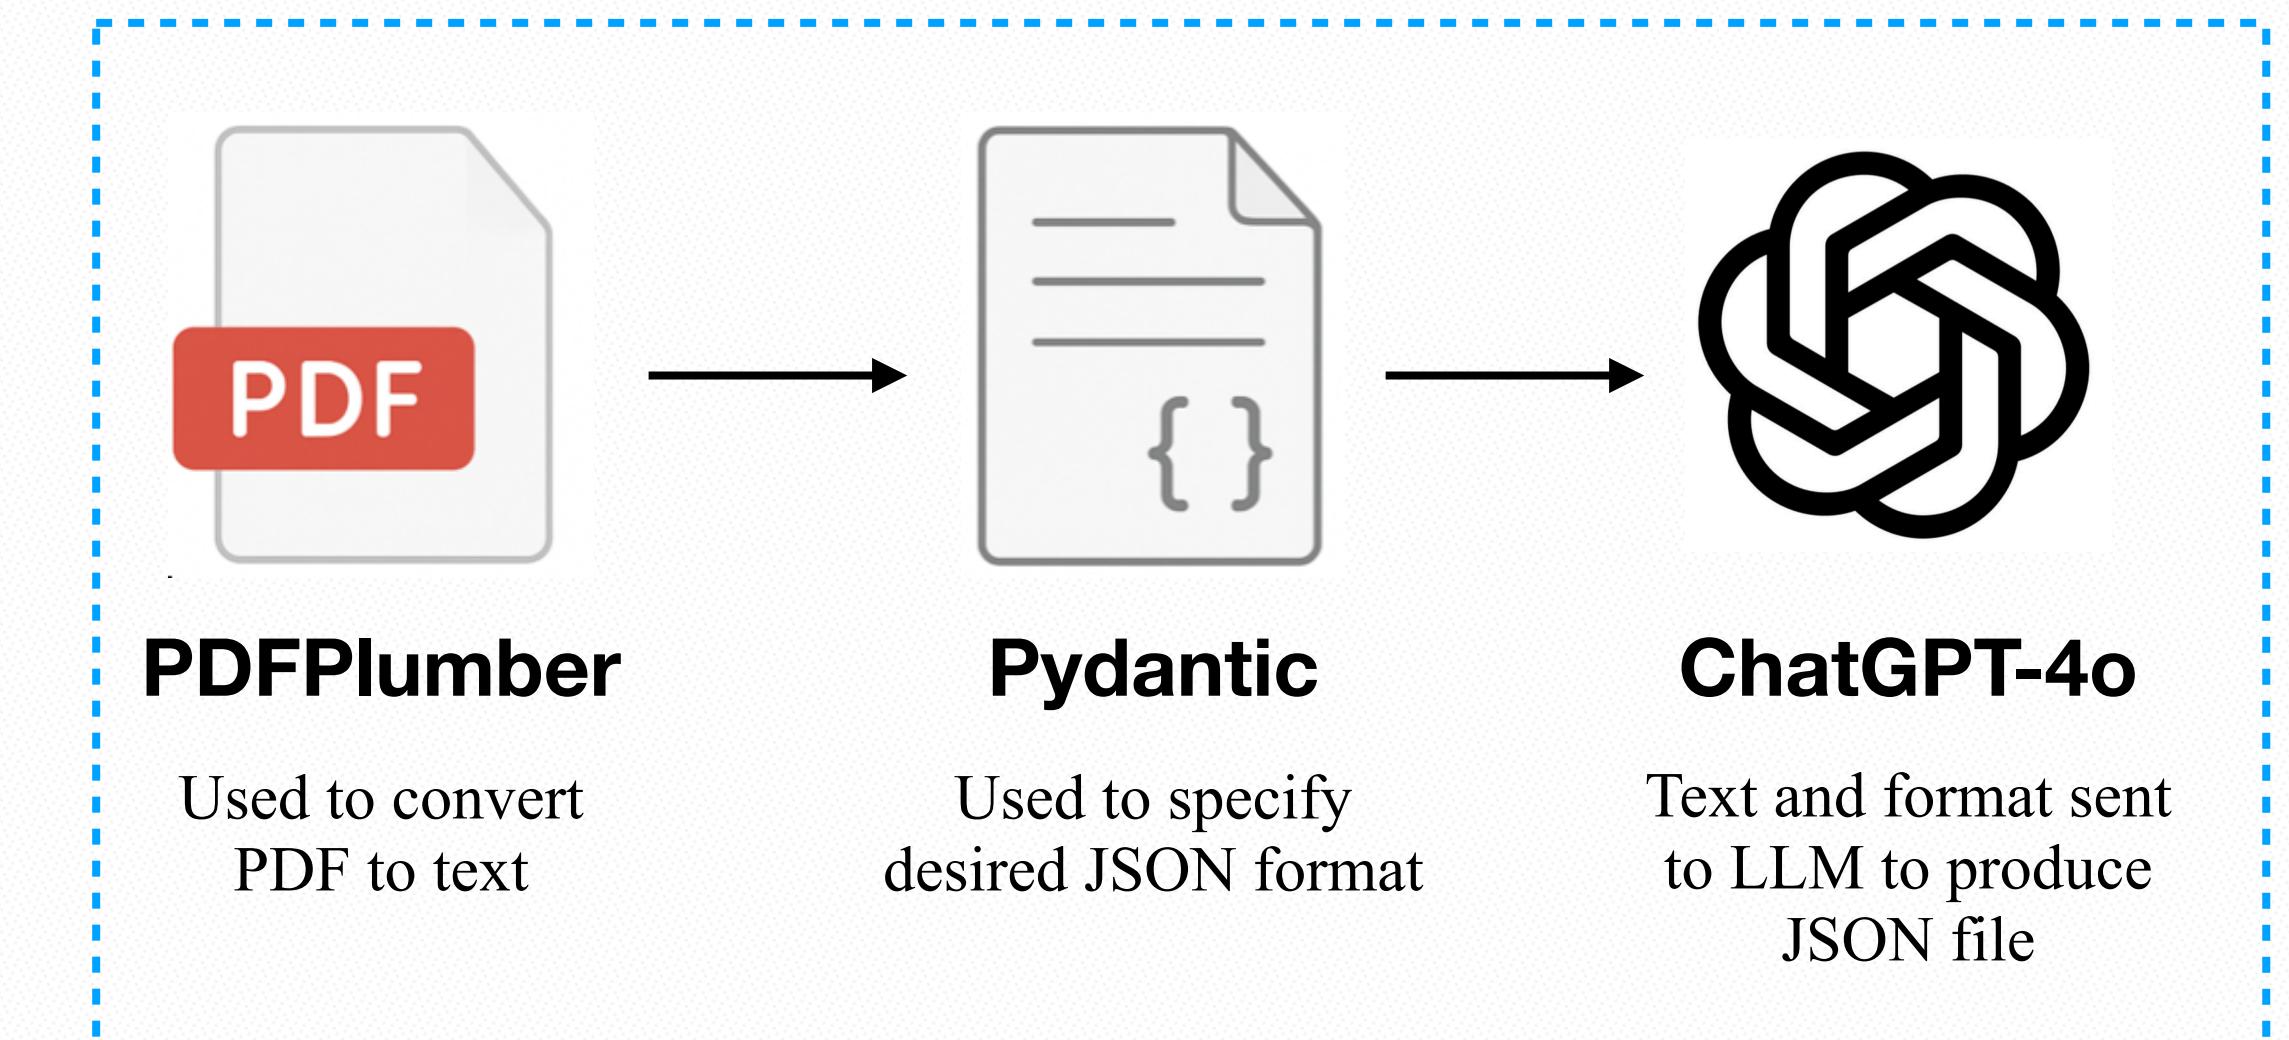
- Used to convert PDF to text

- **Pydantic**

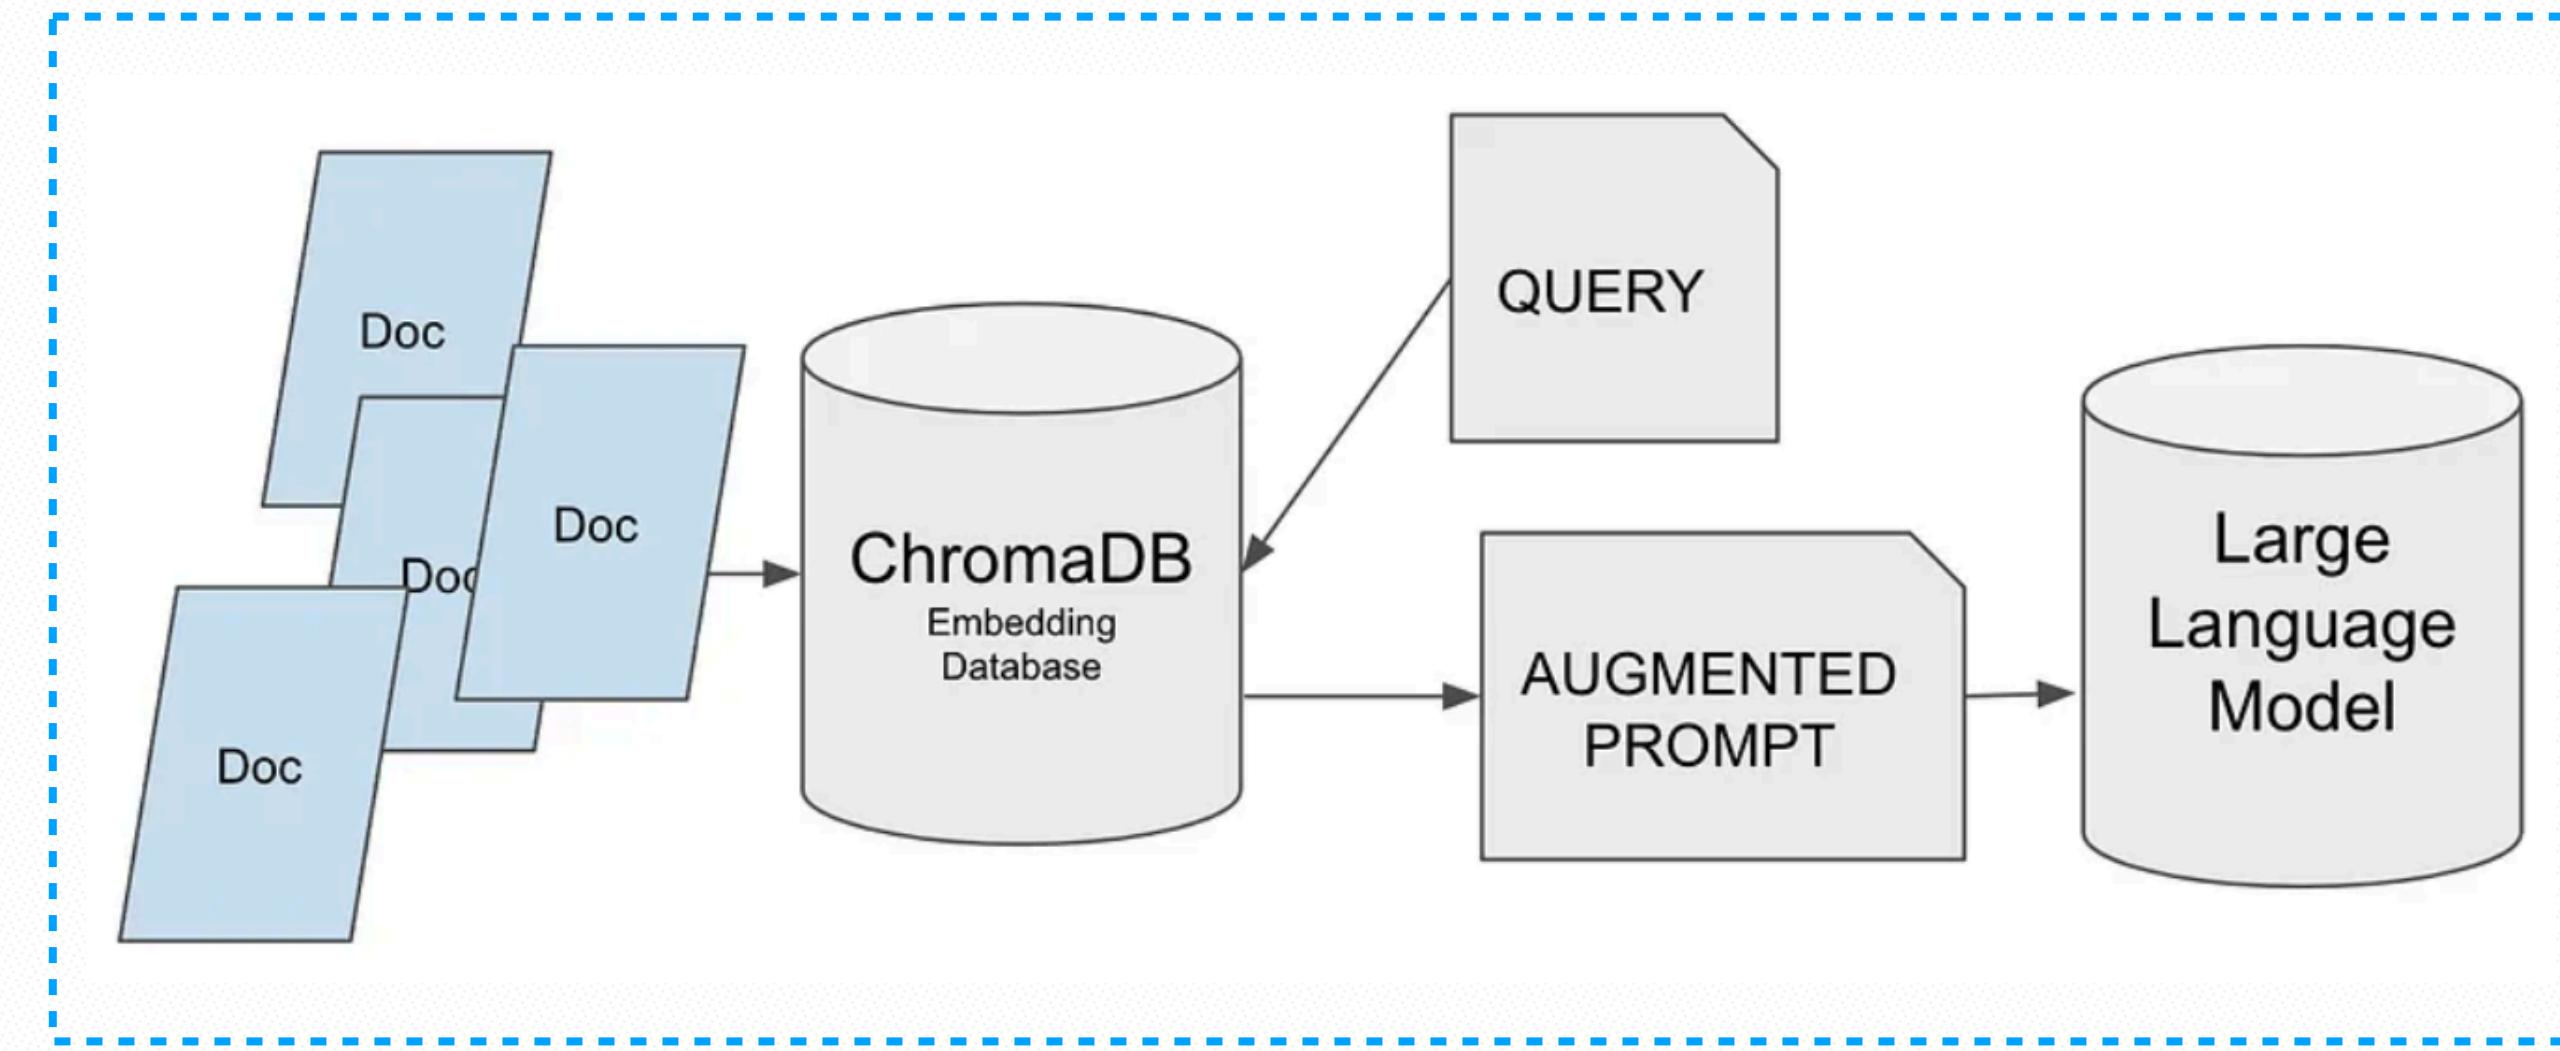
- Used to specify desired JSON format
- Text lines get organized neatly into boxes { key: value }

- **ChatGPT-4o**

- Finally, text and format sent to LLM to produce JSON file
- Neatly formatted JSON file produced by the AI

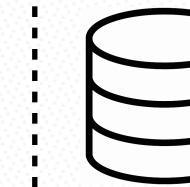


ChromaDB for Retrieval Augmented Generation (RAG)



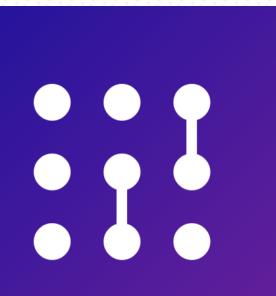
RAG Architecture

- Create a knowledge base for the LLM to reference
- Reduce hallucinations by grounding responses to the retrieved information



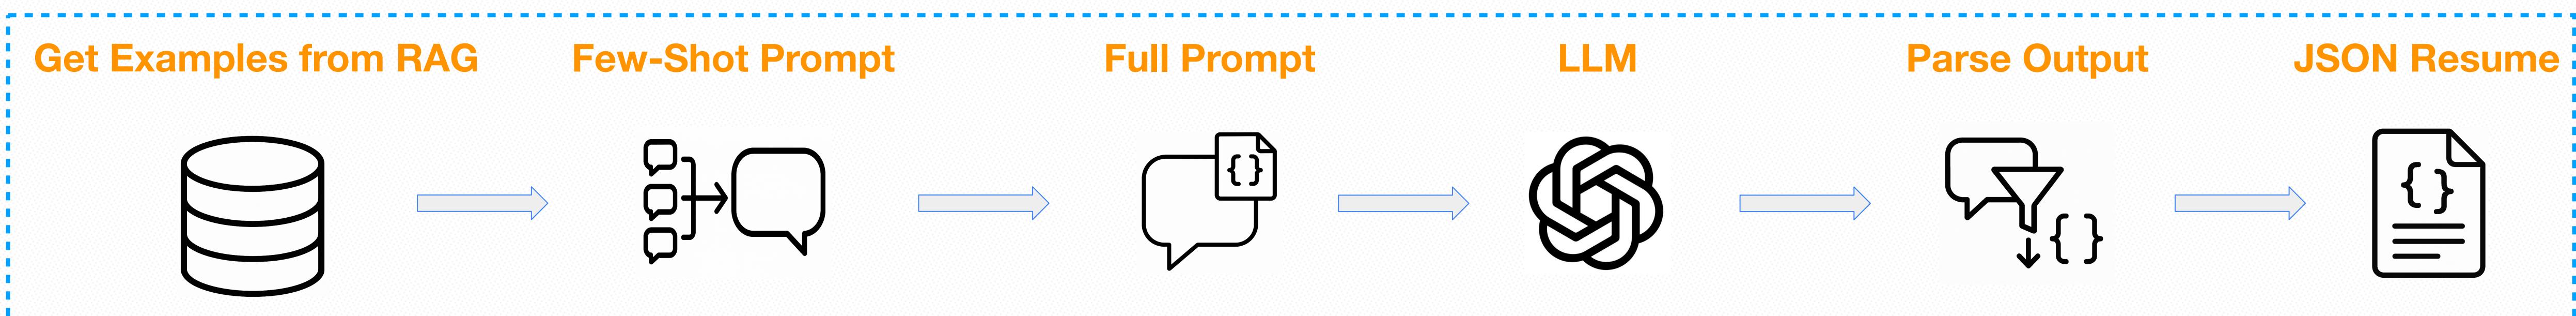
ChromaDB

- Open source vector database
- Stores data as numerical vectors with embeddings
- Quickly queries relevant documents with similarity metrics



GRIDS

Langchain Framework Flow



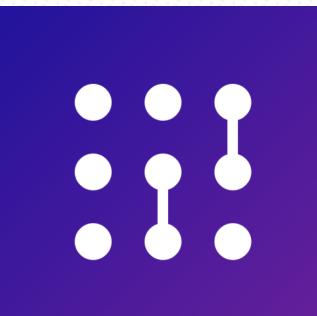
Input Chain: Raw resume is ingested through a Chain with a custom **Prompt Template** to extract the major/field.

Retriever Query: A Retriever searches ChromaDB for **top 2 similar resumes** based on the extracted field.

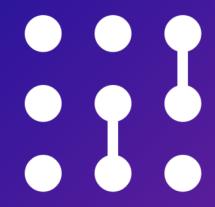
Few-Shot Prompt Assembly: Retrieved example resumes are inserted into a **Few-Shot Prompt Template** to condition the LLM.

Resume Generation Chain: The LLM generates an **ATS-friendly resume** using the few-shot examples and original input.

Output Parsing: An **Output Parser** ensures the generated resume has structured fields.



GRIDS



Graduates Rising in Information and Data Science

Application Walkthrough and Result

Comparing Results

John Doe
johndoe@usc.edu

Los Angeles, CA
+1 (213) 999-9999

EDUCATION

University of Southern California, Master of Science - Computer Science May 2025

Analysis of Algorithms, Deep Learning, Machine Learning for Data Science

University of Pittsburgh, Bachelor of Science - Mathematics and Economics July 2022

Probability & Statistics, Linear Algebra, Applied Econometrics, Game Theory, Calculus 1-3

WORK EXPERIENCE

Mclearn Racing Indianapolis

Data Engineer August 2022 - July 2023

- Handled the receipt, cleaning, and preparation of data from clients using SAS, SQL, and Excel to assist data scientists in building marketing mix models that contributed to improved performance.
- Created a program in SAS that automated the refinement of linear regression models for different customer segments, streamlining routine processes and enhancing team efficiency.

SKILLS

- Programming: SAS (base SAS and Macros), SQL
- Supervised Learning: linear and logistic regressions, decision trees, support vector machines (SVM)
- Unsupervised Learning: k-means clustering, principal component analysis (PCA)
- Data Visualization: Excel, Google Sheets

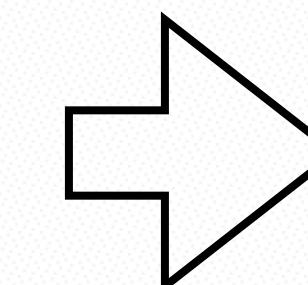
PROJECTS

Fantasy F1 Modeling

- Wanted to stop losing at Fantasy F1 Manager so I aggregated and prepped 5 years of fantasy projection data into a MySQL database
- Built a random forest model in SAS that synthesized the various projections into a unified forecast, leading to improved predictive accuracy compared to other available methods.
- I had heck ton of points on it relative to my previous methodology

Strategy Recommendation Engine

- Aggregated data from F1insight API and used k-nearest-neighbors in SAS to build a better recommendation system for making me predict wins
- I made better pitstop strategies while I play F1 Racing compared my previous methodology, making me win races.



John Doe

johndoe@usc.edu | Los Angeles, CA

SUMMARY

Data scientist with expertise in SAS, SQL, and machine learning algorithms. Experienced in developing predictive models and data visualization techniques.

WORK EXPERIENCE

Data Engineer at Mclearn Racing

2022-08 - 2023-07

Handled the receipt, cleaning, and preparation of data from clients using SAS, SQL, and Excel to assist data scientists in building marketing mix models that contributed to improved performance. Developed and implemented a program in SAS to automate the refinement of linear regression models for diverse customer segments, reducing processing time by 30% and enhancing team productivity. Collaborated with stakeholders to understand data requirements and built pipelines for data ingestion into the data lake.

EDUCATION

University of Southern California - Master of Science in Computer Science

- 2025-05

Analysis of Algorithms, Deep Learning, Machine Learning for Data Science

University of Pittsburgh - Bachelor of Science in Mathematics and Economics

- 2022-07

Probability & Statistics, Linear Algebra, Applied Econometrics, Game Theory, Calculus 1-3

SKILLS

Programming

SAS (base SAS and Macros) SQL Python C#

Machine Learning

Linear and logistic regressions Decision trees Support vector machines (SVM) Predictive analytics

Data Visualization

Power BI Excel Google Sheets

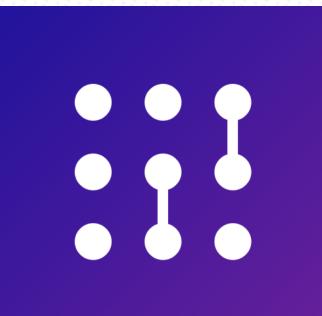
PROJECTS

Fantasy F1 Modeling

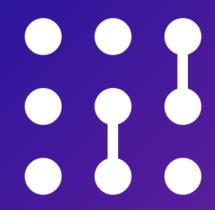
Aggregated and prepped 5 years of fantasy projection data into a MySQL database and built a random forest model in SAS. Enhanced predictive accuracy by 20% compared to existing methods, resulting in a 15% increase in points relative to previous methodology.

Strategy Recommendation Engine

Aggregated data from F1insight API and used k-nearest-neighbors in SAS to build a better recommendation system for predicting wins. Optimized pitstop strategies in F1 Racing, contributing to a 10% increase in race wins.



GRIDS



Graduates Rising in Information and Data Science

THANK YOU

[ress-gen.streamlit.app](#)